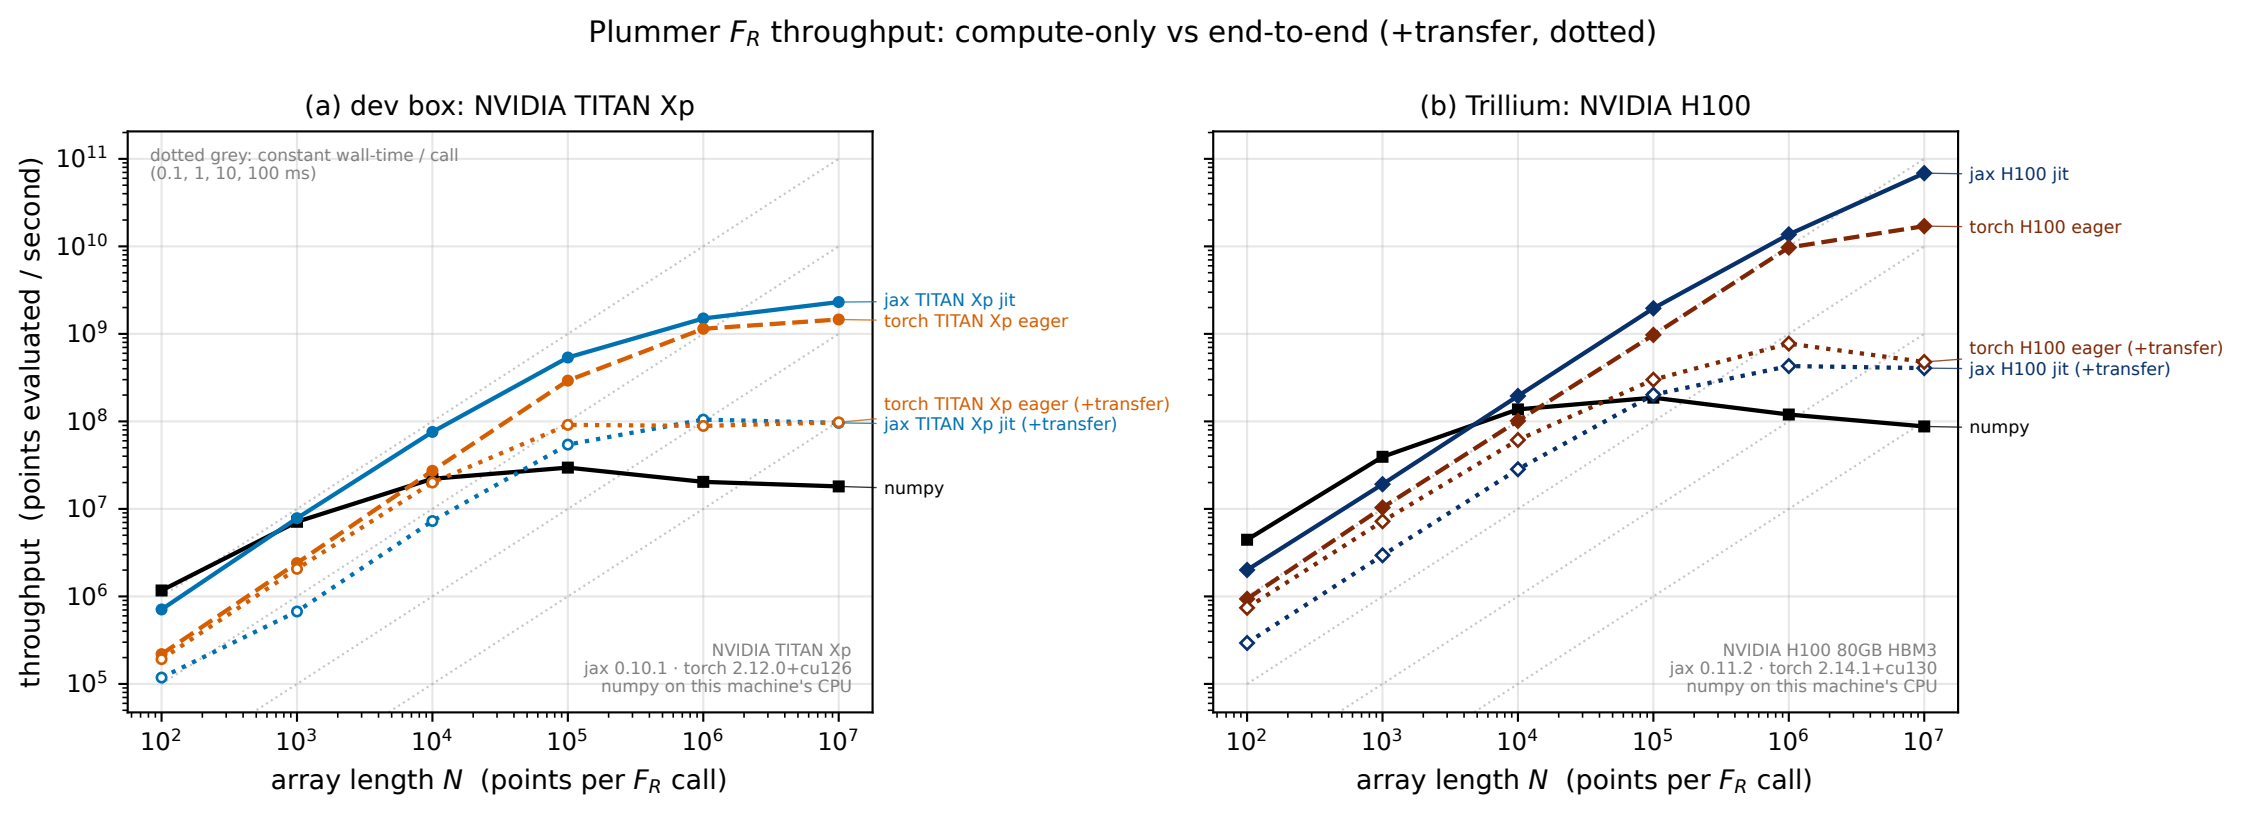
Plummer F_R throughput: compute-only vs end-to-end (+transfer, dotted)
(a) dev box: NVIDIA TITAN Xp
(b) Trillium: NVIDIA H100
throughput (points evaluated / second)
array length N (points per F_R call)
dotted grey: constant wall-time / call
(0.1, 1, 10, 100 ms)
jax TITAN Xp jit
torch TITAN Xp eager
torch TITAN Xp eager (+transfer)
jax TITAN Xp jit (+transfer)
numpy
NVIDIA TITAN Xp
jax 0.10.1 · torch 2.12.0+cu126
numpy on this machine's CPU
jax H100 jit
torch H100 eager
torch H100 eager (+transfer)
jax H100 jit (+transfer)
numpy
NVIDIA H100 80GB HBM3
jax 0.11.2 · torch 2.14.1+cu130
numpy on this machine's CPU