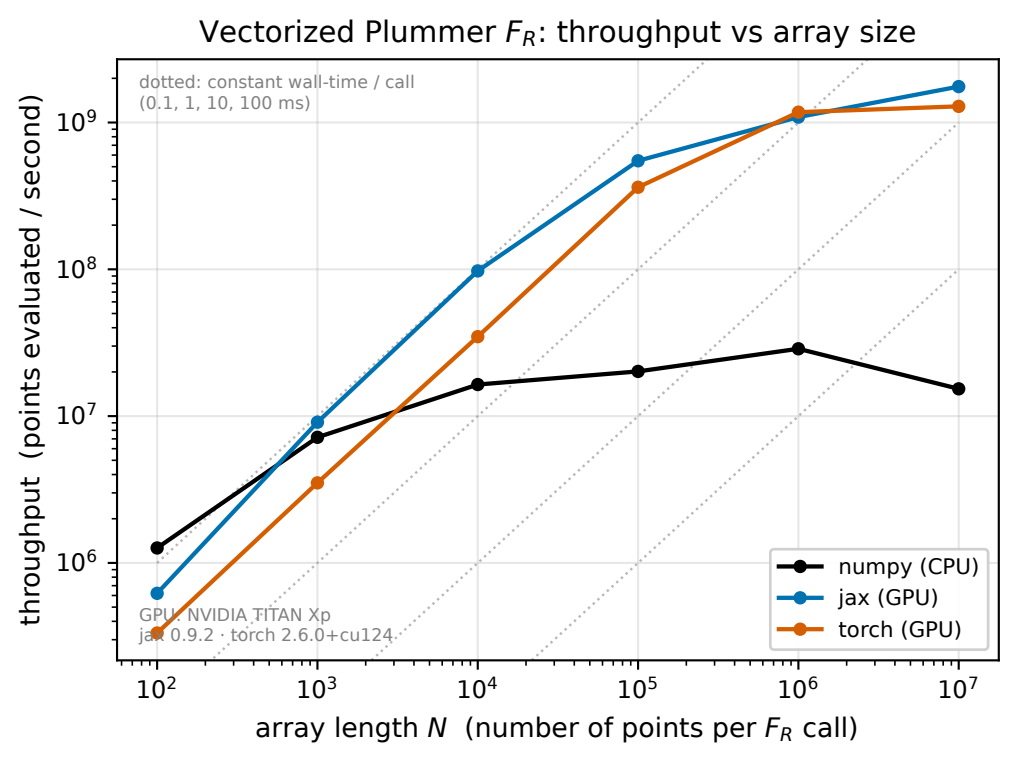
Vectorized Plummer $F_R$: throughput vs array size
dotted: constant wall-time / call
(0.1, 1, 10, 100 ms)
GPU: NVIDIA TITAN Xp
jax 0.9.2 · torch 2.6.0+cu124
numpy (CPU)
jax (GPU)
torch (GPU)
throughput (points evaluated / second)
array length $N$ (number of points per $F_R$ call)
$10^6$
$10^7$
$10^8$
$10^9$
$10^2$
$10^3$
$10^4$
$10^5$
$10^6$
$10^7$